

IoT_SmartParking_XGBoost

February 25, 2025

0.0.1 XGBoost for Predicting Parking Occupancy Status on Feature Engineered Dataset

```
[1]: # Install xgboost if already not installed  
!pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\mahesh\anaconda3\lib\site-packages (2.1.4)  
Requirement already satisfied: numpy in c:\users\mahesh\anaconda3\lib\site-packages (from xgboost) (1.26.4)  
Requirement already satisfied: scipy in c:\users\mahesh\anaconda3\lib\site-packages (from xgboost) (1.11.4)
```

```
[14]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import xgboost as xgb  
from xgboost import XGBClassifier, plot_importance  
from sklearn.preprocessing import MinMaxScaler, LabelEncoder # Import  
↳LabelEncoder here  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import accuracy_score, confusion_matrix, □  
↳classification_report  
import seaborn as sns  
  
# Load the feature engineered IoT Smart Parking dataset  
df = pd.read_csv('IoT_SmartParking_ProCESSED.csv')  
  
# Inspect the data  
print(df.head())  
  
# Preprocessing  
# Assuming 'occupied_spots' is the target variable and other columns are  
↳features  
# The original code used 'occupancy', which was not in the DataFrame  
target_column = 'Occupancy_Status'  
features = [col for col in df.columns if col != target_column and col !=  
↳'Timestamp' and df[col].dtype != object]
```

```

X = df[features].values
y = df[target_column].values

# Normalize features
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

# Encode target variable to numeric using LabelEncoder
encoder = LabelEncoder()
y_encoded = encoder.fit_transform(y)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_encoded,
    test_size=0.2, random_state=42)

# Initialize and train the XGBoost model
model = XGBClassifier(
    objective='binary:logistic', # For binary classification
    colsample_bytree=0.3, # Fraction of features to use for building
trees
    learning_rate=0.1, # Step size at each iteration while moving
towards a minimum
    max_depth=5, # Maximum depth of the trees
    alpha=10, # L1 regularization term on weights
    n_estimators=100, # Number of trees to build
    n_jobs=-1 # Use all CPU threads
)
# Train the model

# Prepare the evaluation data
evals = [(X_train, y_train), (X_test, y_test)]

# Fit the model while tracking evaluation results
evals_result = {}
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

# Get the loss values from evals_result (stored during training)
#vals_result = model.evals_result()

# Extract the logloss for training and validation data
#train_loss = evals_result['validation_0']['logloss']
#val_loss = evals_result['validation_1']['logloss']

```

```

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')

# Print the model parameters (Hyperparameters used for training)
print("Model Parameters:")
print(model.get_params())

# Example if you're using pandas DataFrame
#model.feature_names = X_train.columns.tolist() # Assuming X_train is a pandas DataFrame

print("Model Features:")
print(features)

# Print the number of trees and booster structure (model summary)
# Print the number of trees (using n_estimators)
print(f"\nNumber of trees in the model: {model.get_params()['n_estimators']}")

# Optionally, dump the model structure to a text file
model.get_booster().dump_model('xgboost_model_dump.txt', with_stats=True)

```

	Timestamp	Parking_Spot_ID	Sensor_Reading_Proximity	\
0	2021-01-01 00:00:00.000000000	20	1.023651	
1	2021-01-02 06:39:16.756756756	49	3.903349	
2	2021-01-03 13:18:33.513513513	38	10.315709	
3	2021-01-04 19:57:50.270270270	31	6.588039	
4	2021-01-06 02:37:07.027027027	8	8.213969	

	Sensor_Reading_Pressure	Vehicle_Type_Weight	Vehicle_Type_Height	\
0	1.541461	1831.770127	4.392528	
1	1.621719	1330.815754	4.595638	
2	6.292374	1255.134827	4.313721	
3	1.659870	1523.442919	3.567329	
4	3.278467	1758.490837	5.145836	

	User_Type	Weather_Temperature	Weather_Precipitation	\
0	Visitor	18.092553	1	
1	Registered	13.397533	0	
2	Registered	21.687410	0	
3	Visitor	18.683461	0	
4	Visitor	19.214876	0	

	Nearby_Traffic_Level	...	DayOfWeek_4	DayOfWeek_5	DayOfWeek_6	\
0	Low	...	1.0	0.0	0.0	
1	Low	...	0.0	1.0	0.0	
2	High	...	0.0	0.0	1.0	
3	Medium	...	0.0	0.0	0.0	

```

4           High ...          0.0          0.0          0.0
             DayOfWeek_0.1 DayOfWeek_1.1 DayOfWeek_2.1 DayOfWeek_3.1 DayOfWeek_4.1 \
0            0.0          0.0          0.0          0.0          1.0
1            0.0          0.0          0.0          0.0          0.0
2            0.0          0.0          0.0          0.0          0.0
3            1.0          0.0          0.0          0.0          0.0
4            0.0          0.0          1.0          0.0          0.0

             DayOfWeek_5.1 DayOfWeek_6.1
0            0.0          0.0
1            1.0          0.0
2            0.0          1.0
3            0.0          0.0
4            0.0          0.0

[5 rows x 53 columns]
Accuracy: 100.00%
Model Parameters:
{'objective': 'binary:logistic', 'base_score': None, 'booster': None,
'callbacks': None, 'colsample_bylevel': None, 'colsample_bynode': None,
'colsample_bytree': 0.3, 'device': None, 'early_stopping_rounds': None,
'enable_categorical': False, 'eval_metric': None, 'feature_types': None,
'gamma': None, 'grow_policy': None, 'importance_type': None,
'interaction_constraints': None, 'learning_rate': 0.1, 'max_bin': None,
'max_cat_threshold': None, 'max_cat_to_onehot': None, 'max_delta_step': None,
'max_depth': 5, 'max_leaves': None, 'min_child_weight': None, 'missing': nan,
'monotone_constraints': None, 'multi_strategy': None, 'n_estimators': 100,
'n_jobs': -1, 'num_parallel_tree': None, 'random_state': None, 'reg_alpha':
None, 'reg_lambda': None, 'sampling_method': None, 'scale_pos_weight': None,
'subsample': None, 'tree_method': None, 'validate_parameters': None,
'verbosity': None, 'alpha': 10}
Model Features:
['Parking_Spot_ID', 'Sensor_Reading_Proximity', 'Sensor_Reading_Pressure',
'Vehicle_Type_Weight', 'Vehicle_Type_Height', 'Weather_Temperature',
'Weather_Precipitation', 'Entry_Time', 'Exit_Time', 'Electric_Vehicle',
'Reserved_Status', 'Occupancy_Rate', 'Payment_Amount', 'Parking_Violation',
'Sensor_Reading_Ultrasonic', 'Parking_Duration', 'Environmental_Noise_Level',
'Dynamic_Pricing_Factor', 'Proximity_To_Exit', 'User_Parking_History', 'Hour',
'Month', 'IsWeekend', 'Occupancy_Status_Numeric', 'RollingAvg_Occupancy',
'Prev_Occupancy', 'Prev2_Occupancy', 'Rainfall', 'Temperature',
'Hourly_Occupancy', 'Daily_Occupancy', 'DayOfWeek_0', 'DayOfWeek_1',
'DayOfWeek_2', 'DayOfWeek_3', 'DayOfWeek_4', 'DayOfWeek_5', 'DayOfWeek_6',
'DayOfWeek_0.1', 'DayOfWeek_1.1', 'DayOfWeek_2.1', 'DayOfWeek_3.1',
'DayOfWeek_4.1', 'DayOfWeek_5.1', 'DayOfWeek_6.1']

```

Number of trees in the model: 100

```
[10]: import pandas as pd
import numpy as np
import xgboost as xgb
from xgboost import XGBClassifier, plot_importance
from sklearn.preprocessing import MinMaxScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
import matplotlib.pyplot as plt
import seaborn as sns

# Load dataset (replace 'data.csv' with your actual file)
df = pd.read_csv("IoT_SmartParking_ProCESSED.csv")

# Preprocessing
target_column = 'Occupancy_Status'
features = [col for col in df.columns if col != target_column and col != 'Timestamp' and df[col].dtype != object]

X = df[features].values
y = df[target_column].values

# Normalize features
scaler = MinMaxScaler()
X_scaled = scaler.fit_transform(X)

# Encode target variable to numeric using LabelEncoder
encoder = LabelEncoder()
y_encoded = encoder.fit_transform(y)

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_encoded, test_size=0.2, random_state=42)

# Initialize and train the XGBoost model
model = XGBClassifier(
    objective='binary:logistic', # For binary classification
    colsample_bytree=0.3, # Fraction of features to use for building trees
    learning_rate=0.1, # Step size at each iteration while moving towards a minimum
    max_depth=5, # Maximum depth of the trees
    alpha=10, # L1 regularization term on weights
    n_estimators=100, # Number of trees to build
    n_jobs=-1 # Use all CPU threads
)
```

```

# Prepare the evaluation data
evals = [(X_train, y_train), (X_test, y_test)]

# Fit the model while tracking evaluation results
evals_result = {}
model.fit(X_train, y_train, eval_set=evals, verbose=True)

# Get the predictions
y_pred = model.predict(X_test)

# Get the loss values from evals_result (stored during training)
evals_result = model.evals_result()

# Extract the logloss for training and validation data
train_loss = evals_result['validation_0']['logloss']
val_loss = evals_result['validation_1']['logloss']

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy * 100:.2f}%')

# Print the model parameters
print("Model Parameters:")
print(model.get_params())

```

[0]	validation_0-logloss:0.68531	validation_1-logloss:0.69696
[1]	validation_0-logloss:0.68305	validation_1-logloss:0.69632
[2]	validation_0-logloss:0.67933	validation_1-logloss:0.69447
[3]	validation_0-logloss:0.67674	validation_1-logloss:0.69494
[4]	validation_0-logloss:0.58898	validation_1-logloss:0.60463
[5]	validation_0-logloss:0.58686	validation_1-logloss:0.60376
[6]	validation_0-logloss:0.58368	validation_1-logloss:0.60326
[7]	validation_0-logloss:0.58173	validation_1-logloss:0.60090
[8]	validation_0-logloss:0.57975	validation_1-logloss:0.60043
[9]	validation_0-logloss:0.50915	validation_1-logloss:0.52737
[10]	validation_0-logloss:0.50745	validation_1-logloss:0.52748
[11]	validation_0-logloss:0.44853	validation_1-logloss:0.46636
[12]	validation_0-logloss:0.44801	validation_1-logloss:0.46567
[13]	validation_0-logloss:0.44661	validation_1-logloss:0.46528
[14]	validation_0-logloss:0.39684	validation_1-logloss:0.41356
[15]	validation_0-logloss:0.39607	validation_1-logloss:0.41271
[16]	validation_0-logloss:0.39607	validation_1-logloss:0.41271
[17]	validation_0-logloss:0.39555	validation_1-logloss:0.41193
[18]	validation_0-logloss:0.35302	validation_1-logloss:0.36774
[19]	validation_0-logloss:0.35302	validation_1-logloss:0.36774
[20]	validation_0-logloss:0.31625	validation_1-logloss:0.32953
[21]	validation_0-logloss:0.31625	validation_1-logloss:0.32953
[22]	validation_0-logloss:0.31057	validation_1-logloss:0.32401

[23]	validation_0-logloss:0.30540	validation_1-logloss:0.31978
[24]	validation_0-logloss:0.27478	validation_1-logloss:0.28782
[25]	validation_0-logloss:0.27478	validation_1-logloss:0.28782
[26]	validation_0-logloss:0.24796	validation_1-logloss:0.25981
[27]	validation_0-logloss:0.24796	validation_1-logloss:0.25981
[28]	validation_0-logloss:0.24465	validation_1-logloss:0.25718
[29]	validation_0-logloss:0.24465	validation_1-logloss:0.25718
[30]	validation_0-logloss:0.22144	validation_1-logloss:0.23286
[31]	validation_0-logloss:0.22144	validation_1-logloss:0.23286
[32]	validation_0-logloss:0.22144	validation_1-logloss:0.23286
[33]	validation_0-logloss:0.21941	validation_1-logloss:0.23126
[34]	validation_0-logloss:0.21741	validation_1-logloss:0.23009
[35]	validation_0-logloss:0.21741	validation_1-logloss:0.23009
[36]	validation_0-logloss:0.21566	validation_1-logloss:0.22878
[37]	validation_0-logloss:0.21566	validation_1-logloss:0.22878
[38]	validation_0-logloss:0.21566	validation_1-logloss:0.22878
[39]	validation_0-logloss:0.21566	validation_1-logloss:0.22878
[40]	validation_0-logloss:0.21566	validation_1-logloss:0.22878
[41]	validation_0-logloss:0.21566	validation_1-logloss:0.22878
[42]	validation_0-logloss:0.21425	validation_1-logloss:0.22781
[43]	validation_0-logloss:0.19457	validation_1-logloss:0.20697
[44]	validation_0-logloss:0.19457	validation_1-logloss:0.20697
[45]	validation_0-logloss:0.17713	validation_1-logloss:0.18849
[46]	validation_0-logloss:0.17713	validation_1-logloss:0.18849
[47]	validation_0-logloss:0.17713	validation_1-logloss:0.18849
[48]	validation_0-logloss:0.17713	validation_1-logloss:0.18849
[49]	validation_0-logloss:0.17624	validation_1-logloss:0.18809
[50]	validation_0-logloss:0.17624	validation_1-logloss:0.18809
[51]	validation_0-logloss:0.17624	validation_1-logloss:0.18809
[52]	validation_0-logloss:0.17550	validation_1-logloss:0.18750
[53]	validation_0-logloss:0.17550	validation_1-logloss:0.18750
[54]	validation_0-logloss:0.17550	validation_1-logloss:0.18750
[55]	validation_0-logloss:0.16018	validation_1-logloss:0.17120
[56]	validation_0-logloss:0.14654	validation_1-logloss:0.15667
[57]	validation_0-logloss:0.14654	validation_1-logloss:0.15667
[58]	validation_0-logloss:0.13438	validation_1-logloss:0.14372
[59]	validation_0-logloss:0.13438	validation_1-logloss:0.14372
[60]	validation_0-logloss:0.13438	validation_1-logloss:0.14372
[61]	validation_0-logloss:0.13438	validation_1-logloss:0.14372
[62]	validation_0-logloss:0.13408	validation_1-logloss:0.14353
[63]	validation_0-logloss:0.12325	validation_1-logloss:0.13197
[64]	validation_0-logloss:0.12310	validation_1-logloss:0.13189
[65]	validation_0-logloss:0.11342	validation_1-logloss:0.12156
[66]	validation_0-logloss:0.10475	validation_1-logloss:0.11230
[67]	validation_0-logloss:0.09698	validation_1-logloss:0.10399
[68]	validation_0-logloss:0.09698	validation_1-logloss:0.10399
[69]	validation_0-logloss:0.09698	validation_1-logloss:0.10399
[70]	validation_0-logloss:0.09698	validation_1-logloss:0.10399

```

[71] validation_0-logloss:0.09698 validation_1-logloss:0.10399
[72] validation_0-logloss:0.09698 validation_1-logloss:0.10399
[73] validation_0-logloss:0.09698 validation_1-logloss:0.10399
[74] validation_0-logloss:0.09698 validation_1-logloss:0.10399
[75] validation_0-logloss:0.09000 validation_1-logloss:0.09653
[76] validation_0-logloss:0.09000 validation_1-logloss:0.09653
[77] validation_0-logloss:0.09000 validation_1-logloss:0.09653
[78] validation_0-logloss:0.09000 validation_1-logloss:0.09653
[79] validation_0-logloss:0.09000 validation_1-logloss:0.09653
[80] validation_0-logloss:0.08374 validation_1-logloss:0.08982
[81] validation_0-logloss:0.08374 validation_1-logloss:0.08982
[82] validation_0-logloss:0.08374 validation_1-logloss:0.08982
[83] validation_0-logloss:0.07810 validation_1-logloss:0.08379
[84] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[85] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[86] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[87] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[88] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[89] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[90] validation_0-logloss:0.07303 validation_1-logloss:0.07836
[91] validation_0-logloss:0.06846 validation_1-logloss:0.07347
[92] validation_0-logloss:0.06846 validation_1-logloss:0.07347
[93] validation_0-logloss:0.06846 validation_1-logloss:0.07347
[94] validation_0-logloss:0.06435 validation_1-logloss:0.06906
[95] validation_0-logloss:0.06063 validation_1-logloss:0.06508
[96] validation_0-logloss:0.06063 validation_1-logloss:0.06508
[97] validation_0-logloss:0.06063 validation_1-logloss:0.06508
[98] validation_0-logloss:0.05729 validation_1-logloss:0.06150
[99] validation_0-logloss:0.05427 validation_1-logloss:0.05826


Accuracy: 100.00%
Model Parameters:
{'objective': 'binary:logistic', 'base_score': None, 'booster': None,
'callbacks': None, 'colsample_bylevel': None, 'colsample_bynode': None,
'colsample_bytree': 0.3, 'device': None, 'early_stopping_rounds': None,
'enable_categorical': False, 'eval_metric': None, 'feature_types': None,
'gamma': None, 'grow_policy': None, 'importance_type': None,
'interaction_constraints': None, 'learning_rate': 0.1, 'max_bin': None,
'max_cat_threshold': None, 'max_cat_to_onehot': None, 'max_delta_step': None,
'max_depth': 5, 'max_leaves': None, 'min_child_weight': None, 'missing': nan,
'monotone_constraints': None, 'multi_strategy': None, 'n_estimators': 100,
'n_jobs': -1, 'num_parallel_tree': None, 'random_state': None, 'reg_alpha': None,
'reg_lambda': None, 'sampling_method': None, 'scale_pos_weight': None,
'subsample': None, 'tree_method': None, 'validate_parameters': None,
'verbosity': None, 'alpha': 10}

```


```
[15]: # Print feature importance
plt.figure(figsize=(10, 6))
```

```
plot_importance(model, importance_type='weight', max_num_features=10,  
                 title="Top 10 Important Features")  
plt.show()
```

<Figure size 1000x600 with 0 Axes>



```
[16]: # Plot the training and validation loss  
plt.figure(figsize=(12, 5))  
plt.plot(train_loss, label='Training Loss')  
plt.plot(val_loss, label='Validation Loss')  
plt.xlabel('Boosting Rounds')  
plt.ylabel('Log Loss')  
plt.legend()  
plt.title('Training and Validation Loss')  
plt.show()
```




0.0.2 As can be seen above, the training and validation loss curves are well aligned

0.0.3 and loss exponentially reduced

```
[17]: # Plot the confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 6))
plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.show()

# Print the classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```



Classification Report:


	precision	recall	f1-score	support
0	1.00	1.00	1.00	103
1	1.00	1.00	1.00	97
accuracy			1.00	200
macro avg	1.00	1.00	1.00	200
weighted avg	1.00	1.00	1.00	200

```
[18]: # Plot actual vs predicted values
plt.figure(figsize=(10, 5))
```

```

plt.plot(y_test[:100], label='Actual Occupancy', marker='o')
plt.plot(y_pred[:100], label='Predicted Occupancy', marker='x')
plt.xlabel('Samples')
plt.ylabel('Occupancy')
plt.legend()
plt.title('Actual vs. Predicted Parking Occupancy')
plt.show()

```



0.0.4 As can be seen above F1-Score, Accuracy, Precision and Recall of 100% achieved using XGBoost

0.0.5 for prediction of Parking Occupancy Status using Feature Engineered Dataset